



FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Recurrent Neural Networks

Katharina Breininger, Mingxuan Gu, Noah Maul, Zhaoya Pan, Luca Reeb, Florian Thamm,
Sulaiman Vesal, Tobias Würfl, Zijin Yang
Pattern Recognition Lab, Friedrich-Alexander University of Erlangen-Nürnberg
December 10, 2019






FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Activation Functions



Sigmoid Activation Function




Sigmoid (logistic function)

$$f(x) = \frac{1}{1 + \exp(-x)}$$

$$f'(x) = f(x)(1 - f(x))$$

→ Observe that the derivative can be solely expressed in terms of the activation!

Tanh Activation Function



Tanh

$$f(x) = \tanh(x)$$

$$f'(x) = 1 - f(x)^2$$

→ The derivative is still a function of the activation!



FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Elman Recurrent Neural Network



General strategy

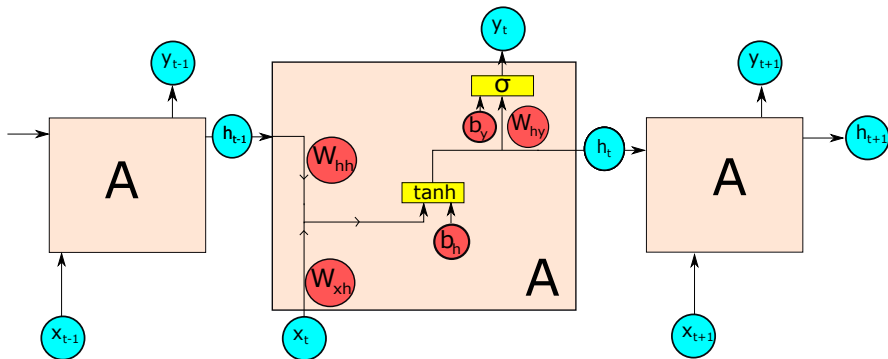
- We interpret the **batch** dimension as **time** dimension now

General strategy

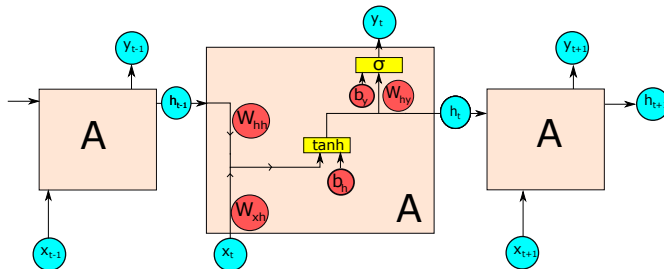
- We interpret the **batch** dimension as **time** dimension now
- Samples are correlated in this dimension

General strategy

- We interpret the **batch** dimension as **time** dimension now
- Samples are correlated in this dimension
- This allows to **reuse** loss functions, optimizers, initializers, activation functions and the Neural Network class



Elman RNN Cell



Output formula:

$$\mathbf{y}_t = \sigma(\mathbf{h}_t \cdot \mathbf{W}_{hy} + \mathbf{b}_y)$$

\mathbf{W}_{hy} : Weight matrix for current hidden state \mathbf{h}_t

\mathbf{b}_h : Output bias

A word on software engineering

- In terms of **encapsulation** - how good was the idea to demand exposition of the weights as member?


A word on software engineering

- In terms of **encapsulation** - how good was the idea to demand exposition of the weights as member?
- Suppose we implement the RNN cell as **composite** structure
- **Getters** and **Setters** provide us the flexibility to do so

A word on software engineering


- In terms of **encapsulation** - how good was the idea to demand exposition of the weights as member?
- Suppose we implement the RNN cell as **composite** structure
- **Getters** and **Setters** provide us the flexibility to do so
- Takeaway? Not doing **proper software engineering** most of the time will demand a price at some point.

Elman RNN Cell



$$\mathbf{h}_t = \tanh(\mathbf{h}_{t-1} \cdot \mathbf{W}_{hh} + \mathbf{x}_t \cdot \mathbf{W}_{xh} + \mathbf{b}_h)$$

Elman RNN Cell



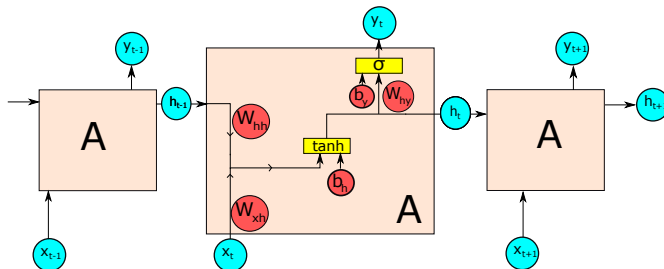
$$\mathbf{h}_t = \tanh(\mathbf{h}_{t-1} \cdot \mathbf{W}_{hh} + \mathbf{x}_t \cdot \mathbf{W}_{xh} + \mathbf{b}_h)$$

\mathbf{W}_{hh} : Weight matrix for previous hidden state \mathbf{h}_{t-1}

\mathbf{W}_{xh} : Weight matrix for current input \mathbf{x}_t

\mathbf{b}_h : Update bias

Elman RNN Cell




$$\mathbf{h}_t = \tanh(\tilde{\mathbf{x}}_t \cdot \mathbf{W}_h)$$


\mathbf{W}_h : Weight matrix of a fully connected layer

$\tilde{\mathbf{x}}_t$: Concatenation of \mathbf{x}_t , \mathbf{h}_{t-1} and a 1

Different from output: Not processed independently!



- Implemented by passing the whole sequence as a **batch**



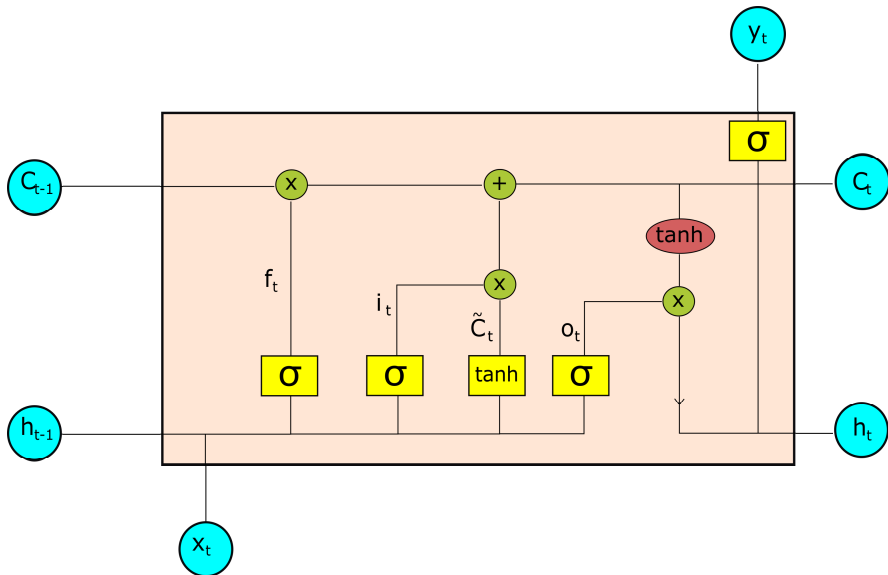
- Implemented by passing **overlapping** parts as a **batch**
- We need to implement memory **Between states**
- Simply store the **last hidden state** and implement a **method** switching whether this state is reused in subsequent forward passes.
- Data has to be fed in **accordingly!**
- Referencing the TBPTT Algorithm presented in the lecture: k_1 is always the sequence length and k_2 is always the TBPTT length.



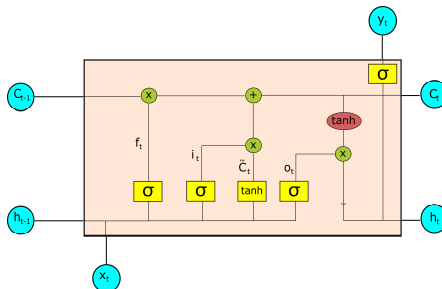
FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Long Short-Term Memory






Forward




- We can reuse a **fully connected** layer again to for the **output**
- The **concatenation** is also **analogous** to the RNN
- The gates σ and the yellow \tanh can be a single **fully connected** layer with an output size of $4 \cdot \text{dim}(\text{hidden state})$
- Remember that we have to pass the vectors of the input tensor **sequentially**

Backward



- Most gradients are again handled by the **embedded layers**
- Again **store and feed the values for backprop externally** to the embedded layers because of **multiple calls to forward**
- We need gradients through **summation, multiplication and copying**

Backward



Sum

Multiply

Copy

$$f(x_1, x_2) = x_1 + x_2$$

$$\frac{\partial \hat{y}}{\partial x_1} = 1$$

Gradient is **copying** $\frac{\partial L}{\partial \hat{y}}$

$$f(x_1, x_2) = x_1 \cdot x_2$$

$$\frac{\partial \hat{y}}{\partial x_1} = x_2$$

Gradient is \cdot with **switched inputs**

Backward pass of sum
So the gradient is a sum!



Thanks for listening.
Any questions?